



Lecture 13 - Strings

Meng-Hsun Tsai
CSIE, NCKU

```
char s1[6] = { 'H', 'e', 'l', 'l', 'o', '\0' };
char s2[6] = "Hello";
char s3[6] = "He" "llo";
```

H	e			o	\0
---	---	--	--	---	----

Introduction

- This lecture covers both **string constants** (or *literals*, as they're called **in the C standard**) and **string variables**.
- Strings are **arrays of characters** in which **a special character—the null character—marks the end**.
- The C library provides a collection of functions for **working with strings**.

String Literals

- A **string literal** is a sequence of characters enclosed within double quotes:
"When you come to a fork in the road, take it."
- String literals may contain escape sequences.
- Character escapes often appear in `printf` and `scanf` format strings.
- For example, each `\n` character in the string
"Candy\nIs dandy\nBut liquor\nIs quicker.\n --Ogden Nash\n"

causes the cursor to advance to the next line:

Candy
Is dandy
But liquor
Is quicker.
--Ogden Nash

Continuing a String Literal

- There are **two ways** to continue a string literal.
- The **backslash character (\)** can be **used to continue a string literal from one line to the next**:

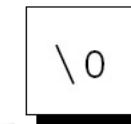
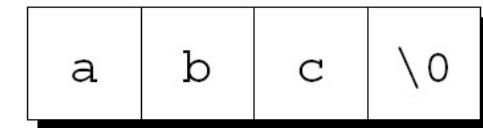
```
printf("When you come to a fork in the road, take it.\n--Yogi Berra");
```

- When **two or more string literals are adjacent**, the compiler will join them into a single string. This rule allows us to **split a string literal over two or more lines**:

```
printf("When you come to a fork in the road, take it.\n"--Yogi Berra");
```

How String Literals Are Stored

- When a C compiler encounters a string literal of length n , it sets aside $n + 1$ bytes of memory for the string.
- This memory will contain the characters in the string, plus one extra character—the **null character**—to mark the end of the string.
- The null character is a byte escape sequence whose bits are all zero, so it's represented by the \0.
- The string literal "abc" is stored as an array of four characters:
- Similarly, the string "" is stored as a single null character:



How String Literals Are Stored (cont.)

- Since a **string literal** is stored as an array, the **compiler treats it as a pointer of type `char *`.**
- Both `printf` and `scanf` **expect a value of type `char *` as their first argument.**
- The following call of `printf` passes the **address of "abc"** (a pointer to where **the letter a** is stored in memory):

```
printf("abc");
```

Operations on String Literals

- We can use a string literal wherever C allows a `char *` pointer:

```
char *p;
```

```
p = "abc";
```

- This assignment makes `p` point to the first character of the string.

Operations on String Literals (cont.)

- String literals **can be subscripted**:

```
char ch;
```

```
ch = "abc"[1];
```

The new value of **ch** will be the letter **b**.

- A function that **converts a number between 0 and 15 into the equivalent hex digit**:

```
char digit_to_hex_char(int digit)
{
    return "0123456789ABCDEF"[digit];
}
```

Operations on String Literals (cont.)

- Attempting to modify a string literal causes undefined behavior:

```
char *p = "abc";  
  
*p = 'd';      / *** WRONG *** /
```

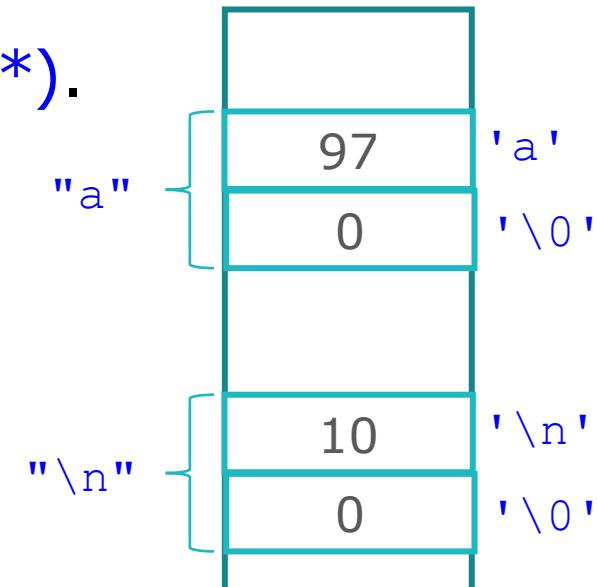
- A program that tries to change a string literal may crash or behave erratically.

String Literals versus Character Constants

- A string literal containing a single character isn't the same as a character constant.
 - "a" is represented by *a pointer (char *)*.
 - 'a' is represented by *an integer*.
- A **legal call** of printf:

```
printf ("\n");
```
- An **illegal call**:

```
printf ('\n');      /*** WRONG *** /
```



String Variables

- Any one-dimensional array of characters can be used to store a string.
- A string must be terminated by a null character.
- Difficulties with this approach:
 - It can be hard to tell whether an array of characters is being used as a string.
 - String-handling functions must be careful to deal properly with the null character.
 - Finding the length of a string requires searching for the null character.

String Variables (cont.)

- If a string variable needs to hold 80 characters, it must be declared to have length 81:

```
#define STR_LEN 80  
...  
char str[STR_LEN+1];
```

- Adding 1 to the desired length allows room for the null character at the end of the string.
- Defining a macro that represents 80 and then adding 1 separately is a common practice.

String Variables (cont.)

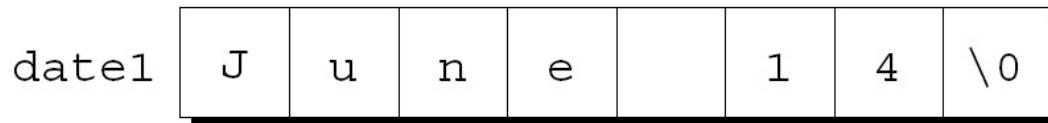
- Be sure to leave room for the null character when declaring a string variable.
- Failing to do so may cause unpredictable results when the program is executed.
- The actual length of a string depends on the position of the terminating null character.
- An array of `STR_LEN + 1` characters can hold strings with lengths between 0 and `STR_LEN`.

Initializing a String Variable

- A string variable **can be initialized** at the same time it's **declared**:

```
char date1[8] = "June 14";
```

- The **compiler will automatically add a null character** so that date1 can be used as a string:

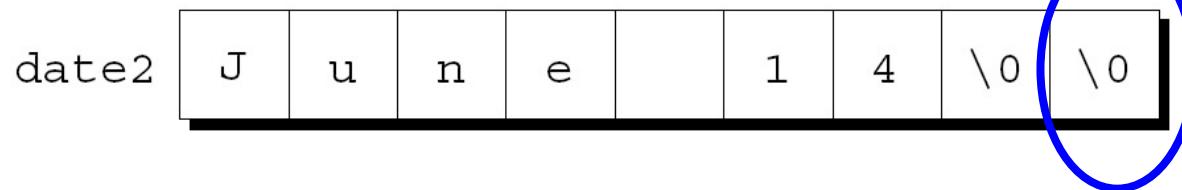


Initializing a String Variable (cont.)

- If the initializer is too short to fill the string variable, the compiler adds extra null characters:

```
char date2[9] = "June 14";
```

Appearance of date2:

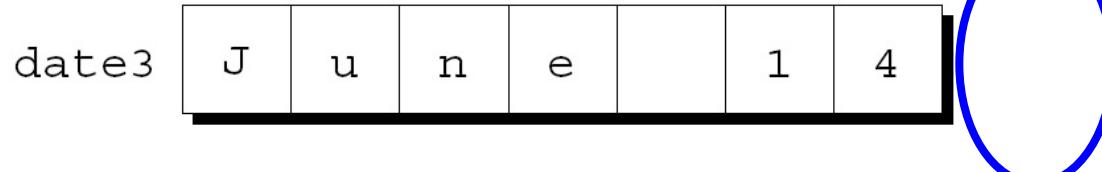


Initializing a String Variable (cont.)

- An **initializer** for a string variable **can't be longer than the variable**, but it **can be the same length**:

```
char date3[7] = "June 14";
```

- There's **no room for the null character**, so the **compiler makes no attempt to store one**:



Initializing a String Variable (cont.)

- The **declaration** of a string variable **may omit its length**, in which case the **compiler computes it**:

```
char date4[] = "June 14";
```

- The **compiler sets aside eight characters** for `date4`, enough to store the characters in "June 14" **plus a null character**.
- Omitting the length of a string variable **is especially useful if the initializer is long**, since computing the length by hand is error-prone.

Character Arrays versus Character Pointers

- The declaration

```
char date[] = "June 14";
```

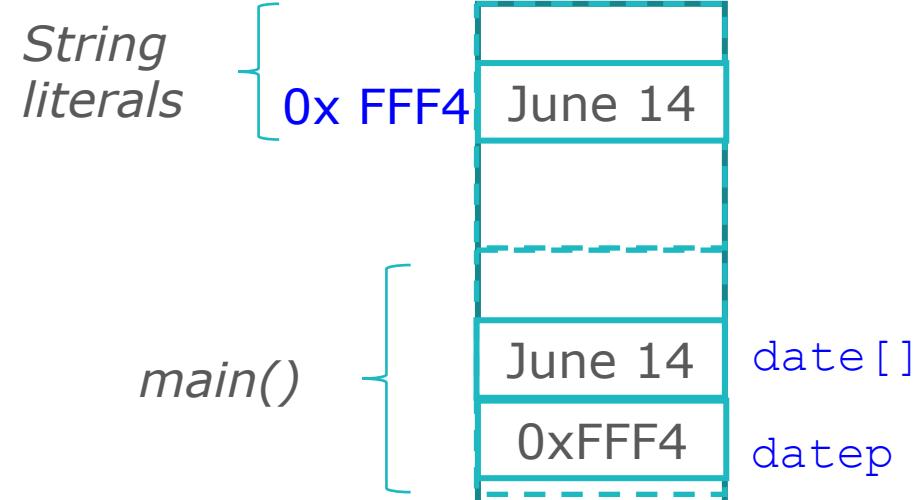
declares `date` to be **an array**,

- The similar-looking

```
char *datep = "June 14";
```

declares `datep` to be **a pointer**.

- Thanks to the close relationship between arrays and pointers, either version can be used as a string.



Character Arrays versus Character Pointers (cont.)

- However, there are **significant differences** between the two versions of date.
 - In the **array version**, the **characters stored in date can be modified**.
In the **pointer version**, datep points to a string literal that **shouldn't be modified**.
 - In the **array version**, date **is an array name**.
In the **pointer version**, datep **is a variable that can point to other strings**.

Character Arrays versus Character Pointers (cont.)

- The declaration

```
char *p;
```

does not allocate space for a string.

- Before we can use `p` as a string, it must point to an array of characters.
- One possibility is to make `p` point to a string variable:

```
char str[STR_LEN+1], *p;  
p = str;
```

- Another possibility is to make `p` point to a dynamically allocated string.

Character Arrays versus Character Pointers (cont.)

- Using an uninitialized pointer variable as a string is a serious error.
- An attempt at building the string "abc":

```
char *p;  
  
p[0] = 'a';      / *** WRONG *** /  
p[1] = 'b';      / *** WRONG *** /  
p[2] = 'c';      / *** WRONG *** /  
p[3] = '\0';     / *** WRONG *** /
```

- Since p hasn't been initialized, this causes undefined behavior.

Reading and Writing Strings

- Writing a string is easy using either `printf` or `puts`.
- Reading a string is a bit harder, because the input may be longer than the string variable into which it's being stored.
- To read a string in a single step, we can use either `scanf` or `gets`.
- As an alternative, we can read strings one character at a time.

Writing Strings Using `printf` and `puts`

- The `%s` conversion specification allows `printf` to write a string:

```
char str[] = "Are we having fun yet?";  
printf("%s\n", str);
```

The output will be

Are we having fun yet?

- `printf` writes the characters in a string one by one until it encounters a null character.

Writing Strings Using `printf` and `puts` (cont.)

- To **print part of a string**, use the conversion specification `%.ps`.
- *p* is the **number of characters** to be displayed.
- The statement

```
printf("%.ps\n", str);
```

will print

Are we

Writing Strings Using `printf` and `puts` (cont.)

- The `%ms` conversion will display a string in a field of size m .
- If the string has fewer than m characters, it will be right-justified within the field.
- To force left justification instead, we can put a minus sign in front of m .
- The m and p values can be used in combination.
- A conversion specification of the form `%m.ps` causes the first p characters of a string to be displayed in a field of size m .

Writing Strings Using `printf` and `puts` (cont.)

- `printf` isn't the only function that can write strings.
- The C library also provides `puts`:
`puts(str);`
- After writing a string, `puts` always writes an additional new-line character.

Reading Strings Using `scanf` and `gets`

- The `%s` conversion specification **allows** `scanf` to read a string into a character array:
`scanf ("%s", str);`
- `str` is treated as a pointer, so there's no need to put the `&` operator in front of `str`.
- When `scanf` is called, it **skips white space**, then **reads characters and stores them in** `str` until it encounters a **white-space character**.
- `scanf` always **stores a null character at the end** of the string.

Reading Strings Using `scanf` and `gets` (cont.)

- `scanf` won't usually read a full line of input.
- A new-line character will cause `scanf` to stop reading, but so will a space or tab character.
- To read an entire line of input, we can use `gets`.
- Properties of `gets`:
 - Doesn't skip white space before starting to read input.
 - Reads until it finds a new-line character.
 - Discards the new-line character instead of storing it; the null character takes its place.

Reading Strings Using `scanf` and `gets` (cont.)

- Consider the following program fragment:

```
char sentence[SENT_LEN+1];  
  
printf("Enter a sentence:\n");  
scanf("%s", sentence);
```

- Suppose that after the prompt

Enter a sentence:

the user enters the line

To C, or not to C: that is the question.

- `scanf` will store the string "To" in `sentence`.

Reading Strings Using `scanf` and `gets` (cont.)

- Suppose that we **replace** `scanf` by `gets`:
`gets (sentence) ;`
- When the user enters the same input as before, **gets will store the string**
" To C, or not to C: that is the question."
in sentence.

Reading Strings Using `scanf` and `gets` (cont.)

- As they read characters into an array, `scanf` and `gets` have no way to detect when it's full.
- Consequently, they may store characters past the end of the array, causing undefined behavior.
- `scanf` can be made safer by using the conversion specification `%ns` instead of `%s`.
- *n* is an integer indicating the maximum number of characters to be stored.
- `gets` is inherently unsafe; `fgets` is a much better alternative.

Reading Strings Character by Character (cont.)

- Programmers often write their own input functions.
- Issues to consider:
 - Should the function skip white space before beginning to store the string?
 - What character causes the function to stop reading: a new-line character, any white-space character, or some other character? Is this character stored in the string or discarded?
 - What should the function do if the input string is too long to store: discard the extra characters or leave them for the next input operation?

Reading Strings Character by Character (cont.)

- Suppose we need a function that (1) doesn't skip white-space characters, (2) stops reading at the first new-line character (which isn't stored in the string), and (3) discards extra characters.
- A prototype for the function:

```
int read_line(char str[], int n);
```

- If the input line contains more than n characters, `read_line` will discard the additional characters.
- `read_line` will return the number of characters it stores in `str`.

Reading Strings Character by Character (cont.)

- `read_line` consists primarily of a loop that calls `getchar` to read a character and then stores the character in `str`, provided that there's room left:

```
int read_line(char str[], int n)
{
    int ch, i = 0;

    while ((ch = getchar()) != '\n')
        if (i < n)
            str[i++] = ch;
    str[i] = '\0'; /* terminates string */
    return i;         /* number of characters stored */
}
```

- `ch` has `int` type rather than `char` type because `getchar` returns an `int` value.

Reading Strings Character by Character (cont.)

- Before returning, `read_line` puts a null character at the end of the string.
- Standard functions such as `scanf` and `gets` automatically put a null character at the end of an input string.
- If we're writing our own input function, we must take on that responsibility.

Accessing the Characters in a String

- Since strings are stored as arrays, we can use subscripting to access the characters in a string.
- To process every character in a string `s`, we can set up a loop that increments a counter `i` and selects characters via the expression `s[i]`.
- A function that counts the number of spaces in a string:

```
int count_spaces(const char s[])
{
    int count = 0, i;
    for (i = 0; s[i] != '\0'; i++)
        if (s[i] == ' ')
            count++;
    return count;
}
```

Accessing the Characters in a String (cont.)

- A version that uses pointer arithmetic instead of array subscripting :

```
int count_spaces(const char *s)
{
    int count = 0;
    for (; *s != '\0'; s++)
        if (*s == ' ')
            count++;
    return count;
}
```

```
int count_spaces(const char s[])
{
    int count = 0, i;
    for (i = 0; s[i] != '\0'; i++)
        if (s[i] == ' ')
            count++;
    return count;
}
```

Accessing the Characters in a String (cont.)

- Questions raised by the `count_spaces` example:
 - *Is it better to use array operations or pointer operations to access the characters in a string?*
-> We can use either or both. Traditionally, C programmers lean toward using pointer operations.
 - *Should a string parameter be declared as an array or as a pointer?*
-> There's no difference between the two.
 - *Does the form of the parameter (`s[]` or `*s`) affect what can be supplied as an argument?*
-> No.

Using the C String Library

- Some programming languages provide operators that can **copy strings, compare strings, concatenate strings, select substrings**, and the like.
- C's operators, in contrast, are essentially useless for working with strings.
- Strings are treated as arrays in C, so they're restricted in the **same ways as arrays**.
- In particular, they can't be copied or compared using operators.

Using the C String Library (cont.)

- Copying a string into a character array using the = operator is not possible:

```
char str1[10], str2[10];  
...  
str1 = "abc";    /*** WRONG ***/  
str2 = str1;    /*** WRONG ***/
```

Using an array name as the left operand of = is illegal.

- Initializing a character array using = is legal, though:

```
char str1[10] = "abc";
```

In this context, = is not the assignment operator.

Using the C String Library (cont.)

- Attempting to **compare strings** using a **relational or equality operator** **is legal but won't produce the desired result**:

```
if (str1 == str2) ...    / *** WRONG *** /
```

- This statement **compares str1 and str2 as pointers**.
- Since str1 and str2 have different addresses, the expression str1 == str2 must have the value 0.

Using the C String Library (cont.)

- The **C library provides a rich set of functions for performing operations on strings.**
- Programs that need string operations should contain the following line:
`#include <string.h>`
- In subsequent examples, assume that `str1` and `str2` are character arrays used as strings.

The **strcpy** (String Copy) Function

- **Prototype** for the `strcpy` function:

```
char *strcpy(char *s1, const char *s2);
```

- `strcpy` **copies** the **string s2** **into** the **string s1**.
 - To be precise, we should say “`strcpy` copies the string pointed to by `s2` into the array pointed to by `s1`.”
- `strcpy` **returns** `s1` (a pointer to the destination string).

The `strcpy` (String Copy) Function (cont.)

- A call of `strcpy` that stores the string "abcd" in `str2`:

```
strcpy(str2, "abcd");
/* str2 now contains "abcd" */
```

- A call that copies the contents of `str2` into `str1`:

```
strcpy(str1, str2);
/* str1 now contains "abcd" */
```

- In the call `strcpy(str1, str2)`, `strcpy` has no way to check that the `str2` string will fit in the array pointed to by `str1`.
 - If it doesn't, undefined behavior occurs.

The **strncpy** (String Copy) Function

- Calling the **strncpy** function is a safer, albeit slower, way to copy a string.
- **strncpy has a third argument** that limits the number of characters that will be copied.
- **strncpy will leave str1 without a terminating null character** if the length of **str2** is greater than or equal to the size of the **str1** array.
- A safer way to use **strncpy**:

```
strncpy(str1, str2, sizeof(str1) - 1);  
str1[sizeof(str1)-1] = '\0';
```
- The second statement guarantees that **str1** is always null-terminated.

The `strlen` (String Length) Function

- Prototype for the `strlen` function:

```
size_t strlen(const char *s);
```

- `size_t` is a **typedef name** that represents one of C's **unsigned integer types**.
- `strlen` returns the length of a string `s`, not including the null character.
- Examples:

```
int len;
```

```
len = strlen("abc"); /* len is now 3 */
```

```
len = strlen(""); /* len is now 0 */
```

```
strcpy(str1, "abc");
```

```
len = strlen(str1); /* len is now 3 */
```

The **strcat** (String Concatenation) Function

- Prototype for the **strcat** function:

```
char *strcat(char *s1, const char *s2);
```

- **strcat appends the contents of the string s2 to the end of the string s1.**
- It **returns s1** (a pointer to the **resulting string**).
- **strcat examples:**

```
strcpy(str1, "abc");
strcat(str1, "def");
/* str1 now contains "abcdef" */
strcpy(str1, "abc");
strcpy(str2, "def");
strcat(str1, str2);
/* str1 now contains "abcdef" */
```

The **strcat** (String Concatenation) Function (cont.)

- As with **strcpy**, the value returned by **strcat** is normally discarded.
- The following example shows how the return value might be used:

```
strcpy(str1, "abc");
strcpy(str2, "def");
strcat(str1, strcat(str2, "ghi"));
/* str1 now contains "abcdefghi";
   str2 contains "defghi" */
```

The `strcat` (String Concatenation) Function (cont.)

- `strcat(str1, str2)` causes **undefined behavior** if the **str1 array isn't long enough** to accommodate the characters from `str2`.
- Example:

```
char str1[6] = "abc";
```

```
strcat(str1, "def");    /*** WRONG ***/
```

- `str1` is limited to six characters, causing `strcat` to write past the end of the array.

The `strncat` (String Concatenation) Function

- The `strncat` function is a **safer but slower** version of `strcat`.
- Like `strncpy`, it has a **third argument** that limits the number of characters it will copy.
- A call of `strncat`:

```
strncat(str1, str2, sizeof(str1) - strlen(str1) - 1);
```
- `strncat` will terminate `str1` with a null character, which isn't included in the third argument.

The `strcmp` (String Comparison) Function

- Prototype for the `strcmp` function:

```
int strcmp(const char *s1, const char *s2);
```

- `strcmp` compares the strings `s1` and `s2`, returning a value less than, equal to, or greater than 0, depending on whether `s1` is less than, equal to, or greater than `s2`.

- Testing whether `str1` is less than `str2`:

```
if (strcmp(str1, str2) < 0) /* is str1 < str2? */
```

- Testing whether `str1` is less than or equal to `str2`:

```
if (strcmp(str1, str2) <= 0) /* is str1 <= str2? */
```

- By choosing the proper operator (`<`, `<=`, `>`, `>=`, `==`, `!=`), we can test any possible relationship between `str1` and `str2`.

The strcmp (String Comparison) Function (cont.)

- As it compares two strings, **strcmp** looks at the numerical codes for the characters in the strings.
- Important properties of ASCII:
 - A-Z, a-z, and 0-9 have consecutive codes.
 - All upper-case letters are less than all lower-case letters.
 - Digits are less than letters.
 - Spaces are less than all printing characters.

ASCII value	Character	ASCII value	Character	ASCII value	Character
000	^@	043	+	086	V
001	^A	044	,	087	W
002	^B	045	-	088	X
003	^C	046	.	089	Y
004	^D	047	/	090	Z
005	^E	048	0	091	[
006	^F	049	1	092	\
007	^G	050	2	093]
008	^H	051	3	094	,
009	^I	052	4	095	-
010	^J	053	5	096	'
011	^K	054	6	097	a
012	^L	055	7	098	b
013	^M	056	8	099	c
014	^N	057	9	100	d
015	^O	158	:	101	e
016	^P	059	;	102	f
017	^Q	060	<	103	g
018	^R	061	=	104	h
019	^S	062	>	105	i
020	^T	063	?	106	j
021	^U	064	@	107	k
022	^V	065	A	108	l
023	^W	066	B	109	m
024	^X	067	C	110	n
025	^Y	068	D	111	o
026	^Z	069	E	112	p
027	^[_	070	F	113	q
028	^`	071	G	114	r
029	^]	072	H	115	s
030	^~	073	I	116	t
031	^_	074	J	117	u
032	[space]	075	K	118	v
033	!	076	L	119	w
034	"	077	M	120	x
035	#	078	N	121	y
036	\$	079	O	122	z
037	%	080	P	123	{
038	&	081	Q	124	
039	,	082	R	125	}
040	(083	S	126	-
041)	084	T	127	DEL
042	*	085	U		

The `strcmp` (String Comparison) Function (cont.)

- `strcmp` considers `s1` to be less than `s2` if either one of the following conditions is satisfied:
 - The first i characters of `s1` and `s2` match, but the $(i+1)$ st character of `s1` is less than the $(i+1)$ st character of `s2`.
`"abc" < "abd"`
 - All characters of `s1` match `s2`, but `s1` is shorter than `s2`.
`"abc" < "abcd"`

Program: Printing a One-Month Reminder List

- The `remind.c` program prints a **one-month list of daily reminders**.
- The **user will enter a series of reminders**, with each **prefixed by a day** of the month.
- When the **user enters 0** instead of a valid day, the program will print a **list of all reminders entered, sorted by day**.
- The next slide shows a session with the program.

Program: Printing a One-Month Reminder List (cont.)

```
Enter day and reminder: 24 Susan's birthday
Enter day and reminder: 5 6:00 - Dinner with Marge and Russ
Enter day and reminder: 26 Movie - "Chinatown"
Enter day and reminder: 7 10:30 - Dental appointment
Enter day and reminder: 12 Movie - "Dazed and Confused"
Enter day and reminder: 5 Saturday class
Enter day and reminder: 12 Saturday class
Enter day and reminder: 0
```

Day Reminder

```
5 Saturday class
5 6:00 - Dinner with Marge and Russ
7 10:30 - Dental appointment
12 Saturday class
12 Movie - "Dazed and Confused"
24 Susan's birthday
26 Movie - "Chinatown"
```

Program: Printing a One-Month Reminder List (cont.)

- Overall strategy:
 - **Read** a series of day-and-reminder combinations.
 - **Store them in order (sorted by day)**.
 - **Display** them.
- `scanf` will be used to **read the days**.
- `read_line` will be used to **read the reminders**.

Program: Printing a One-Month Reminder List (cont.)

- The **strings** will be stored in a two-dimensional array of characters.
- Each row of the array contains one string.
- Actions taken after the program reads a day and its associated reminder:
 - Search the array to determine where the day belongs, using `strcmp` to do comparisons.
 - Use `strcpy` to move all strings below that point down one position.
 - Copy the day into the array and call `strcat` to append the reminder to the day.

Program: Printing a One-Month Reminder List (cont.)

- One complication: how to right-justify the days in a two-character field.
- A solution: use `scanf` to read the day into an integer variable, than call `sprintf` to convert the day back into string form.
- `sprintf` is similar to `printf`, except that it writes output into a string.
- The following call writes the value of `day` into `day_str`.

```
sprintf(day_str, "%2d", day);
```

- The following call of `scanf` ensures that the user doesn't enter more than two digits:

```
scanf ("%2d", &day);
```

Program: Printing a One-Month Reminder List (cont.)

```
#include <stdio.h>          remind.c
#include <string.h>
#define MAX_REMIND 50      /* maximum number of reminders */
#define MSG_LEN 60         /* max length of reminder message */
int read_line(char str[], int n);
int main(void)
{
    char reminders[MAX_REMIND][MSG_LEN+3];
    char day_str[3], msg_str[MSG_LEN+1];
    int day, i, j, num_remind = 0;

    for (;;) {
        if (num_remind == MAX_REMIND) {
            printf("-- No space left --\n");
            break;
        }
        printf("Enter day and reminder: ");
        scanf("%2d", &day);
        if (day == 0)
            break;
        sprintf(day_str, "%2d", day);
        read_line(msg_str, MSG_LEN);
```

Program: Printing a One-Month Reminder List (cont.)

```
for (i = 0; i < num_remind; i++)
    if (strcmp(day_str, reminders[i]) < 0)
        break;
for (j = num_remind; j > i; j--)
    strcpy(reminders[j], reminders[j-1]);

strcpy(reminders[i], day_str);
strcat(reminders[i], msg_str);
num_remind++;
}

printf("\nDay Reminder\n");
for (i = 0; i < num_remind; i++)
    printf(" %s\n", reminders[i]);
return 0;
}
```

```
int read_line(char str[], int n)
{
    int ch, i = 0;

    while ((ch = getchar()) != '\n')
        if (i < n)
            str[i++] = ch;
    str[i] = '\0';
    return i;
}
```

Arrays of Strings

- There is more than one way to store an array of strings.
- One option is to use a two-dimensional array of characters, with one string per row:

```
char planets[][] = {"Mercury", "Venus", "Earth",
                    "Mars", "Jupiter", "Saturn",
                    "Uranus", "Neptune", "Pluto"};
```

- The number of rows in the array can be omitted, but we must specify the number of columns.

Arrays of Strings (cont.)

- Unfortunately, the planets array contains a fair bit of wasted space (extra null characters):

	0	1	2	3	4	5	6	7
0	M	e	r	c	u	r	y	\0
1	V	e	n	u	s	\0	\0	\0
2	E	a	r	t	h	\0	\0	\0
3	M	a	r	s	\0	\0	\0	\0
4	J	u	p	i	t	e	r	\0
5	S	a	t	u	r	n	\0	\0
6	U	r	a	n	u	s	\0	\0
7	N	e	p	t	u	n	e	\0
8	P	l	u	t	\0	\0	\0	\0

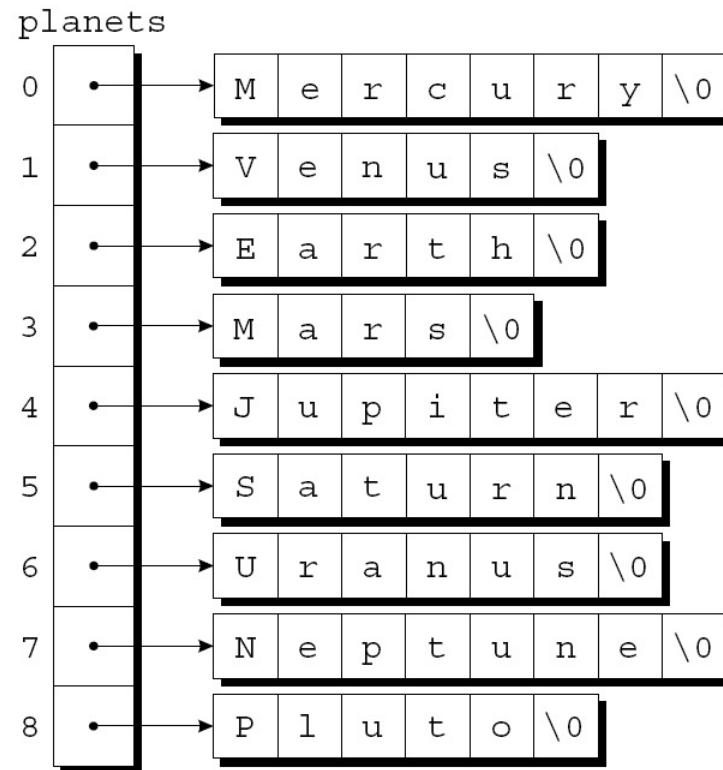
Arrays of Strings (cont.)

- Most collections of strings will have a mixture of long strings and short strings.
- What we need is a **ragged array**, whose rows can have different lengths.
- We can simulate a ragged array in C by creating an array whose elements are *pointers* to strings:

```
char *planets[] = {"Mercury", "Venus", "Earth",
                    "Mars", "Jupiter", "Saturn",
                    "Uranus", "Neptune", "Pluto"};
```

Arrays of Strings (cont.)

- This small change has a dramatic effect on how planets is stored:



Arrays of Strings (cont.)

- To access one of the planet names, all we need do is subscript the planets array.
- Accessing a character in a planet name is done in the same way as accessing an element of a two-dimensional array.
- A loop that searches the planets array for strings beginning with the letter M:

```
for (i = 0; i < 9; i++)
    if (planets[i][0] == 'M')
        printf("%s begins with M\n", planets[i]);
```

Command-Line Arguments

- When we run a program, we'll often need to supply it with information.
- This may include a file name or a switch that modifies the program's behavior.
- Examples of the UNIX ls command:

```
ls
```

```
ls -l
```

```
ls -l remind.c
```

Command-Line Arguments (cont.)

- Command-line information is available to all programs, not just operating system commands.
- To obtain access to ***command-line arguments***, main must have two parameters:

```
int main(int argc, char *argv[])
{
    ...
}
```

- Command-line arguments are called ***program parameters*** in the C standard.

Command-Line Arguments (cont.)

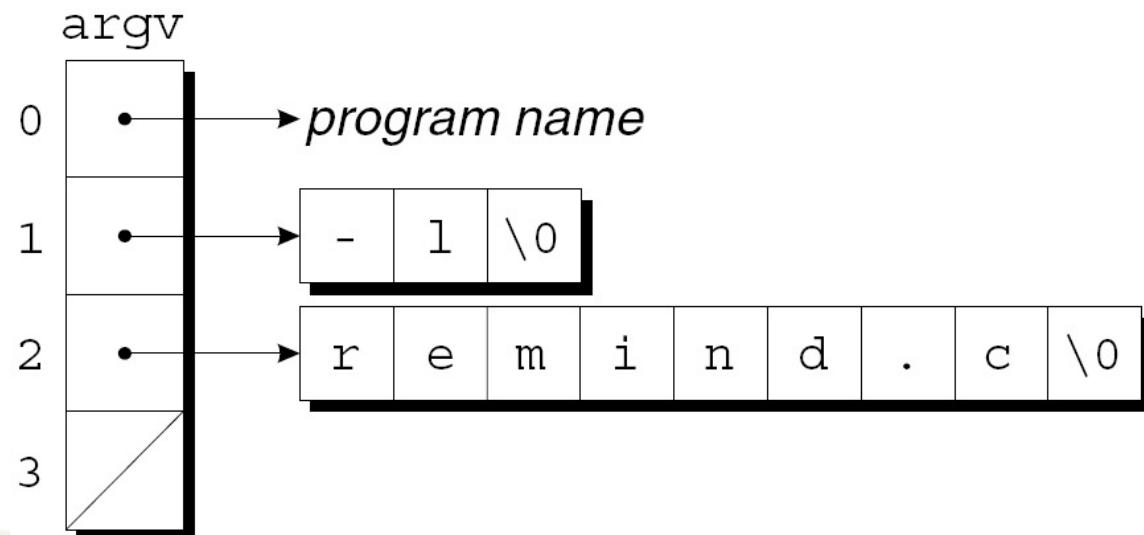
- `argc` ("argument **count**") is **the number of command-line arguments**.
- `argv` ("argument **vector**") is an **array of pointers** to the command-line arguments (stored as strings).
- `argv[0]` points to the **name of the program**, while `argv[1]` through `argv[argc-1]` point to the **remaining command-line arguments**.
- `argv[argc]` is **always a *null pointer***—a special pointer that points to nothing.
 - The macro `NULL` **represents a null pointer**.

Command-Line Arguments (cont.)

- If the user enters the command line

```
ls -l remind.c
```

then `argc` will be 3, and `argv` will have the following appearance:



Command-Line Arguments (cont.)

- Typically, a program that expects command-line arguments will set up a loop that examines each argument in turn.
- One way to write such a loop is to use an integer variable as an index into the `argv` array:

```
int i;
for (i = 1; i < argc; i++)
    printf("%s\n", argv[i]);
```

- Another technique is to set up a pointer to `argv[1]`, then increment the pointer repeatedly:

```
char **p;
for (p = &argv[1]; *p != NULL; p++)
    printf("%s\n", *p);
```

Program: Checking Planet Names

- The `planet.c` program illustrates how to access command-line arguments.
- The program is designed to check a series of strings to see which ones are names of planets.
- The strings are put on the command line:

planet Jupiter venus Earth fred

- The program will indicate whether each string is a planet name and, if it is, display the planet's number:

Jupiter is planet 5

venus is not a planet

Earth is planet 3

fred is not a planet

Program: Checking Planet Names (cont.)

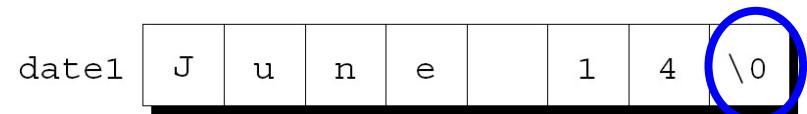
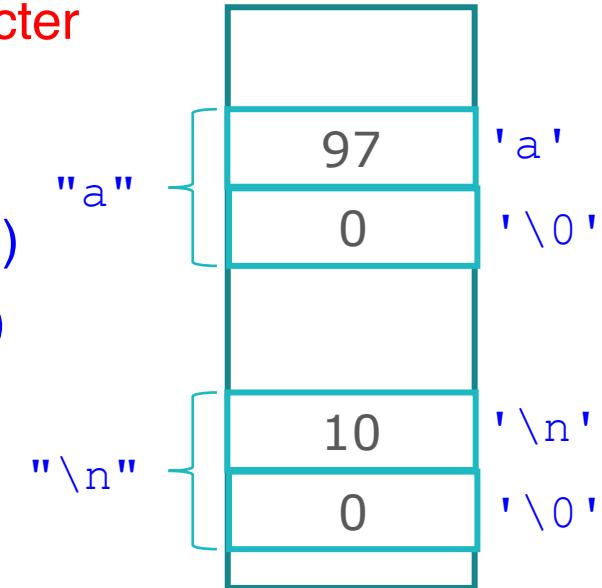
planet.c

```
#include <stdio.h>
#include <string.h>
#define NUM_PLANETS 9
int main(int argc, char *argv[])
{
    char *planets[] = {"Mercury", "Venus", "Earth",
                       "Mars", "Jupiter", "Saturn",
                       "Uranus", "Neptune", "Pluto"};
    int i, j;
    for (i = 1; i < argc; i++) {
        for (j = 0; j < NUM_PLANETS; j++)
            if (strcmp(argv[i], planets[j]) == 0) {
                printf("%s is planet %d\n", argv[i], j + 1);
                break;
            }
        if (j == NUM_PLANETS)
            printf("%s is not a planet\n", argv[i]);
    }
    return 0;
}
```



A Quick Review to This Lecture

- Strings are arrays of characters in which the null character marks the end (an extra byte is required).
- String Literals
 - String Literals: "a" (represented by a *pointer (char *)*)
 - Character Constants: 'a' (represented by an *integer*)
 - String literals can be subscripted:
`char ch = "abc"[1];`
- String Variables:
 - Be sure to leave room for the null character
`char date1[8] = "June 14";`
 - Length can be omitted
`char date1[] = "June 14";`

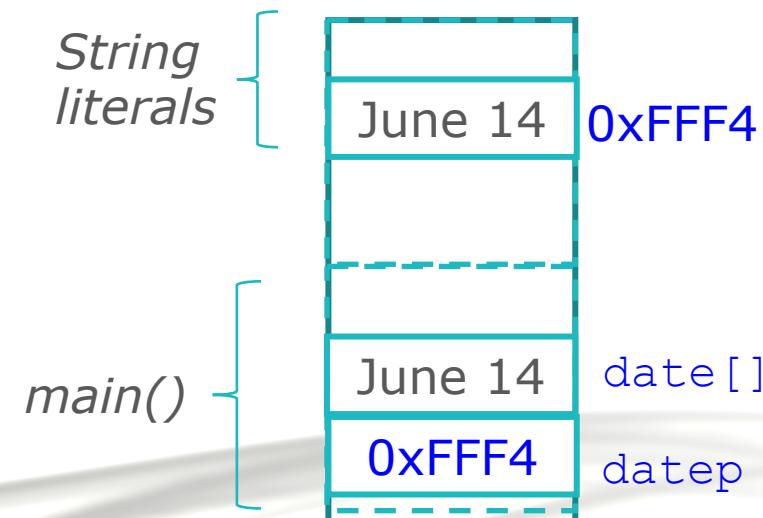


A Quick Review to This Lecture (cont.)

- Character array vs. character pointer

Character array	Character pointer
<code>char date[] = "June 14";</code>	<code>char *datep = "June 14";</code>
Allocate space	Does not allocate space
Characters can be modified	Characters can not be modified

- Writing strings: `printf / puts`
- Reading strings: `scanf / gets / fgets`
- Reading a character: `getchar`



A Quick Review to This Lecture (cont.)

- `printf("%m.ps", str);`
 - `%m.ps`: display **first p characters** in a **field of size m** (right justified)
 - `%-m.ps`: display **first p characters** in a **field of size m** (left justified)
- `puts(str); // always writes an additional new-line character`

A Quick Review to This Lecture (cont.)

- `scanf ("%ns", str);`
 - $%ns$: n indicates the maximum number of stored characters
 - skips leading white spaces, **reads**, and **stops reading** while **encountering a white space** or n characters are read
- `gets(str);`
 - Do not skip leading white spaces, **reads**, and **stops reading** only when encountering a new-line character
- `fgets(str, n, stdin);`

A Quick Review to This Lecture (cont.)

- C String Library (`#include <string.h>`)

- `strcpy / strncpy` copies the string `s2` into the string `s1`, and returns `s1`

```
char *strcpy(char *s1, const char *s2);    strcpy(str2, "abcd");  
char *strncpy(char *s1, const char *s2, size_t n);  
                                         strncpy(str1, str2, sizeof(str1) - 1);  
                                         str1[sizeof(str1)-1] = '\0';
```

- `strlen` returns the length of a string `s`, not including the null character.

```
size_t strlen(const char *s);           len = strlen("abc"); // 3
```

- `strcat / strncat` appends the string `s2` to the end of string `s1`, and returns `s1`

```
char *strcat(char *s1, const char *s2);  strcat(str1, "def");  
char *strncat(char *s1, const char *s2, size_t n);  
                                         strncat(str1, str2, sizeof(str1) - strlen(str1) - 1);
```

A Quick Review to This Lecture (cont.)

- `strcmp` compares the strings `s1` and `s2`, returning a negative value, zero or a positive value if `s1` is less than, equal to, or greater than `s2`.

```
int strcmp(const char *s1, const char *s2);
```

```
if (strcmp(str1, str2) < 0) /* is str1 < str2? */
```

`s1` is less than `s2` if either one of the following conditions is satisfied:

- The first *i* characters of `s1` and `s2` match, but the (*i*+1)st character of `s1` is less than the (*i*+1)st character of `s2`.
`"abc" < "abd"`
- All characters of `s1` match `s2`, but `s1` is shorter than `s2`.
`"abc" < "abcd"`

- `sprintf` is similar to `printf`, except that it writes output into a string.

```
sprintf(day_str, "%2d", day);
```

A Quick Review to This Lecture (cont.)

- To obtain access to **command-line arguments**, `main` must have two parameters:

```
int main(int argc, char *argv[])
```

- If the user enters the command line

```
ls -l remind.c
```

then `argc` will be 3, and `argv` will have the following appearance:

